Exp. No : 5

HIVE Installation

1. Download and Extract apache hive 3.1.2

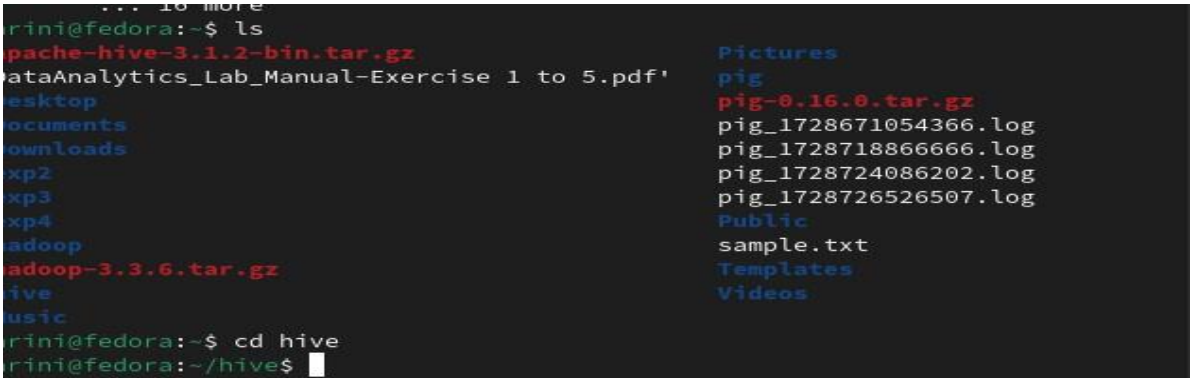

2. Update HIVE Configurations in .bashrc

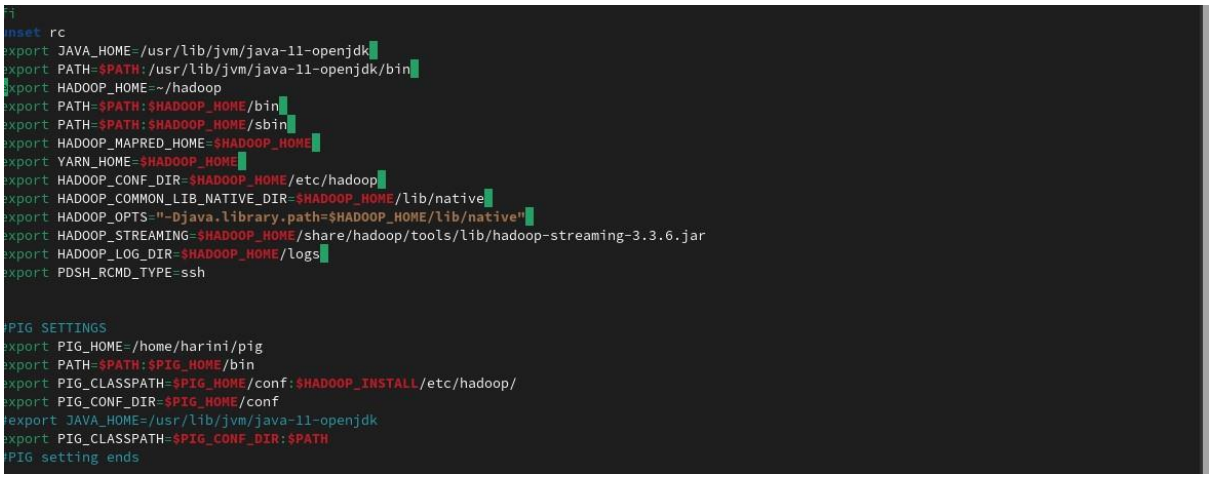

3. Change directory to apache-hive-3.1.2-bin/conf

4. Create hive-env.sh

```
  GNU nano 7.2                                                                hi
#   if [ -z "$DEBUG" ]; then
#     export HADOOP_OPTS="$HADOOP_OPTS -XX:NewRatio=12 -Xms10m -XX:MaxHeapFreeRati
#   else
#     export HADOOP_OPTS="$HADOOP_OPTS -XX:NewRatio=12 -Xms10m -XX:MaxHeapFreeRati
#   fi
# fi

# The heap size of the jvm stared by hive shell script can be controlled via:
#
# export HADOOP_HEAPSIZE=1024
#
# Larger heap size may be required when running queries over large number of files
# By default hive shell scripts use a heap size of 256 (MB).  Larger heap size wou
# appropriate for hive server.


# Set HADOOP_HOME to point to a specific hadoop install directory
# HADOOP_HOME=${bin}/../../hadoop

# Hive Configuration Directory can be controlled by:
# export HIVE_CONF_DIR=

# Folder containing extra libraries required for hive compilation/execution can be
# export HIVE_AUX_JARS_PATH=
```

5. Install and Change mysql root password

```
Enter password:
Welcome to the MySQL monitor.  Commands end with ; or \g.
Your MySQL connection id is 8
Server version: 8.0.39 Source distribution

Copyright (c) 2000, 2024, Oracle and/or its affiliates.

Oracle is a registered trademark of Oracle Corporation and/or its
affiliates. Other names may be trademarks of their respective
owners.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.

mysql> SELECT user, host, plugin FROM mysql.user WHERE user = 'root';
+------+-----------+-----------------------+
| user | host      | plugin                |
+------+-----------+-----------------------+
| root | localhost | mysql_native_password |
+------+-----------+-----------------------+
1 row in set (0.00 sec)
```

6. Create and configure hive-site.xml

```
  </description>
</property>
<property>
  <name>hive.repl.partitions.dump.parallelism</name>
  <value>100</value>
  <description>Number of threads that will be used to dump partition data information during repl dump.</description>
</property>
<property>
  <name>hive.repl.dumpdir.clean.freq</name>
  <value>0s</value>
  <description>
    Expects a time value with unit (d/day, h/hour, m/min, s/sec, ms/msec, us/usec, ns/nsec), which is sec if not specified.
    Frequency at which timer task runs to purge expired dump dirs.
  </description>
</property>
<property>
  <name>hive.repl.dumpdir.ttl</name>
  <value>7d</value>
  <description>
    Expects a time value with unit (d/day, h/hour, m/min, s/sec, ms/msec, us/usec, ns/nsec), which is day if not specified.
    TTL of dump dirs before cleanup.
  </description>
</property>
<property>
  <name>hive.repl.dump.metadata.only</name>
```

7. Download and Move mysql java connector to apache-hive-3.1.2-bin/lib

```
keerthi@fedora:~/hive/conf$ cd
keerthi@fedora:~$ cd hive
keerthi@fedora:~/hive$ cd lib
keerthi@fedora:~/hive/lib$ ls
accumulo-core-1.7.3.jar                          hbase-server-2.0.0-alpha4.jar              jersey-media-jaxb-2.25.1.jar
accumulo-fate-1.7.3.jar                          hbase-shaded-miscellaneous-1.0.1.jar       jersey-server-2.25.1.jar
accumulo-start-1.7.3.jar                         hbase-shaded-netty-1.0.1.jar               jettison-1.1.jar
accumulo-trace-1.7.3.jar                         hbase-shaded-protobuf-1.0.1.jar            jetty-annotations-9.3.20.v20170531.jar
aircompressor-0.10.jar                           HikariCP-2.6.1.jar                         jetty-client-9.3.20.v20170531.jar
ant-1.9.1.jar                                    hive-accumulo-handler-3.1.2.jar            jetty-http-9.3.20.v20170531.jar
ant-launcher-1.9.1.jar                           hive-beeline-3.1.2.jar                     jetty-io-9.3.20.v20170531.jar
antlr4-runtime-4.5.jar                           hive-classification-3.1.2.jar              jetty-jaas-9.3.20.v20170531.jar
antlr-runtime-3.5.2.jar                          hive-cli-3.1.2.jar                         jetty-jndi-9.3.20.v20170531.jar
aopalliance-repackaged-2.5.0-b32.jar             hive-common-3.1.2.jar                      jetty-plus-9.3.20.v20170531.jar
apache-curator-2.12.0.pom                        hive-contrib-3.1.2.jar                     jetty-rewrite-9.3.20.v20170531.jar
apache-jsp-9.3.20.v20170531.jar                  hive-druid-handler-3.1.2.jar               jetty-runner-9.3.20.v20170531.jar
apache-jstl-9.3.20.v20170531.jar                 hive-exec-3.1.2.jar                        jetty-schemas-3.1.jar
arrow-format-0.8.0.jar                           hive-hbase-handler-3.1.2.jar               jetty-security-9.3.20.v20170531.jar
arrow-memory-0.8.0.jar                           hive-hcatalog-core-3.1.2.jar               jetty-server-9.3.20.v20170531.jar
arrow-vector-0.8.0.jar                           hive-hcatalog-server-extensions-3.1.2.jar  jetty-servlet-9.3.20.v20170531.jar
asm-5.0.1.jar                                    hive-hplsql-3.1.2.jar                      jetty-util-9.3.20.v20170531.jar
asm-commons-5.0.1.jar                            hive-jdbc-3.1.2.jar                        jetty-webapp-9.3.20.v20170531.jar
asm-tree-5.0.1.jar                               hive-jdbc-handler-3.1.2.jar                jetty-xml-9.3.20.v20170531.jar
audience-annotations-0.5.0.jar                   hive-kryo-registrator-3.1.2.jar            jline-2.12.jar
avatica-1.11.0.jar                               hive-llap-client-3.1.2.jar                 joda-time-2.9.9.jar
avro-1.7.7.jar                                   hive-llap-common-3.1.2.jar                 joni-2.1.11.jar
bonecp-0.8.0.RELEASE.jar                         hive-llap-common-3.1.2-tests.jar           jpam-1.1.jar
calcite-core-1.16.0.jar                          hive-llap-ext-client-3.1.2.jar             jsch-0.1.54.jar
```

8. Execute schematool -initSchema -dbType mysql

```
lf4j-reload4j-1.7.36.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]
Metastore connection URL:        jdbc:mysql://localhost/metastore?createDatabaseIfNo
tExist=true
Metastore Connection Driver :    com.mysql.cj.jdbc.Driver
Metastore connection User:       root
Starting metastore schema initialization to 3.1.0
Initialization script hive-schema-3.1.0.mysql.sql
```

9. Start hive

```
keerthi@fedora:~/hive/lib$ hive
which: no hbase in (/home/keerthi/.local/bin:/home/keerthi/bin:/usr/local/bin:/usr/local/sbin:/usr/bin:/usr/sbin:/usr/lib/jvm/java-8-open
jdk/bin:/home/keerthi/hadoop/bin:/home/keerthi/hadoop/sbin:/home/keerthi/pig/bin:/home/keerthi/hive/bin:/usr/lib/jvm/java-8-openjdk/bin:/
home/keerthi/hadoop/bin:/home/keerthi/hadoop/sbin:/home/keerthi/pig/bin:/home/keerthi/hive/bin)
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/home/keerthi/hive/lib/log4j-slf4j-impl-2.10.0.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/home/keerthi/hadoop/share/hadoop/common/lib/slf4j-reload4j-1.7.36.jar!/org/slf4j/impl/StaticLoggerBind
er.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.apache.logging.slf4j.Log4jLoggerFactory]
Hive Session ID = ad4ea647-2583-4543-9dd9-9b1264c68925

Logging initialized using configuration in jar:file:/home/keerthi/hive/lib/hive-common-3.1.2.jar!/hive-log4j2.properties Async: true
Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spa
rk, tez) or using Hive 1.X releases.
Hive Session ID = 2943ba5b-4d25-4a33-8009-4bd5e56c6a9d
hive>
```